

**Data Pipeline Integration & Music
Recommendation System on AWS with Docker**

Submitted in partial fulfilment of the requirements for the degree of

**Master of Technology in Data Engineering
Master of Technology in Artificial Intelligence**

by

Abhishek Sahu

(M23CSA504)

Sudarshan P (G23AI1046)

Under the guidance of

Asst Prof. Sumit Kalra

IIT Jodhpur

**Indian Institute of Technology Jodhpur
Software & Data Engineering
MTech. Third Sem (SEP 2024)**



॥ त्वं ज्ञानमयो विज्ञानमयोऽसि ॥

Declaration

I hereby declare that the work presented in this Project Report titled Data Pipeline Integration & Music Recommendation System on AWS + Docker submitted to the Indian Institute of Technology Jodhpur in partial fulfilment of the requirements for the award of the degree of Master of Technology in Data Engineering, is a bonafide record of the research work carried under the supervision of Assistant Professor. Sumit Kalra, IIT - Jodhpur. The contents of this Project Report in full or in parts, have not been submitted to, and will not be submitted by me to, any other Institute or University in India or abroad for the award of any degree or diploma.

Acknowledgement

I would like to express my sincere appreciation and gratitude to all those who contributed to the completion of this report. Their support and assistance were invaluable, and without their contributions, this work would not have been possible.

I am deeply thankful to Assistant Professor. Sumit Kalra, IIT - Jodhpur, for their guidance and mentorship throughout the project. Their expertise and insights greatly enriched the content and quality of this report.

I would also like to extend my thanks to my teammate for their collaboration and assistance in various stages of the research and writing process. Their dedication and hard work played a significant role in shaping the final outcome.

Furthermore, I want to acknowledge the support of IIT Jodhpur, which provided the necessary resources and environment for the successful completion of this project.

Lastly, my heartfelt thanks go to my family and friends for their understanding, encouragement, and patience during the demanding phases of this endeavor.

Abstract

This project presents a comprehensive data pipeline integration and music recommendation system on Amazon Web Services (AWS). The project leverages the Spotify Web API to extract music data, including user playlists and tracks, and integrates it with AWS services to build a scalable and efficient data pipeline.

This project presents a comprehensive approach to integrating Docker with Amazon Web Services (AWS) for deploying a scalable music recommendation system. The project leverages the Spotify Web API to extract music data, including user playlists and tracks, and integrates it with AWS services to build a scalable and efficient data pipeline.

The data pipeline consists of several key components: data ingestion, data transformation, data loading, and data analysis. The data ingestion component uses Dockerized Python scripts to extract data from the Spotify API and store it in an Amazon S3 bucket. The data transformation component, also containerized using Docker, preprocesses the data—handling missing values, normalizing features, and performing other transformations—using AWS Glue. The transformed data is loaded into Amazon S3 for querying and analysis with Amazon Athena.

The music recommendation system applies collaborative filtering algorithms to generate personalized song recommendations based on user listening history and preferences. User-song interactions, such as listening history, preferences, and ratings, are taken into account to deliver highly accurate and personalized recommendations.

This project demonstrates a comprehensive strategy for deploying scalable and efficient systems using Docker and AWS services, including ECR, ECS, S3, Glue, and Athena. The use of Docker for containerization ensures portability, while AWS services provide the necessary infrastructure for scaling and handling large volumes of data, making the solution suitable for production environments.

Introduction & Background

Motivation

The motivation behind this project is to build a scalable and efficient data pipeline that can handle large volumes of music data and provide personalized song recommendations to users.

Overview

This project aims to construct a robust ETL (Extract, Transform, Load) pipeline using Docker and AWS services. The pipeline facilitates seamless extraction, transformation, and loading of data from the Spotify Web API, while Docker containers ensure portability and consistency across environments. In addition to data processing, the project integrates a music recommendation system that utilizes collaborative filtering techniques to deliver personalized song recommendations.

The project architecture combines the benefits of Docker containerization with AWS services to ensure efficient data ingestion, transformation, and analysis. AWS services such as S3, Glue, and Athena provide the necessary infrastructure to handle large volumes of data. The final deliverable will include insightful analytics and visualizations to better understand user behaviors and preferences, enabling more accurate and personalized recommendations for users.

Objectives

- Extract: Retrieve the required data from the Spotify API using Dockerized Python scripts.
- Transform: Process and transform the raw data using Docker containers for consistency & AWS Glue.
- Load: Store the transformed data in Amazon S3 and make it available for querying & analysis using AWS Athena.
- Automation: Using AWS services like CloudWatch for triggering, Glue for transformation Athena for querying.
- Deployment: Ensure a smooth CI/CD process by containerizing the application with Docker and automating testing and deployment pipelines using AWS services.

AWS Services Used

- S3 (Simple Storage Service): To store raw and transformed data.
- AWS Lambda: To run code in response to events for data extraction and transformation.
- CloudWatch: To monitor AWS resources and set alarms for Lambda function execution.
- Glue Crawler: To crawl data in S3 and create a Glue Data Catalog.
- Glue Data Catalog: To manage metadata about the data stored in S3.
- Amazon Athena: To query the data stored in S3 using standard SQL.

Scope

The project involves extracting data from the Spotify API using Dockerized Python scripts, ensuring consistency across environments. The data is transformed using Docker containers and AWS Glue for enrichment and schema management, and stored in Amazon S3 for querying via Athena. A CI/CD pipeline is set up using AWS services like CodePipeline, ECS, and ECR to automate deployment and testing of the ETL pipeline.

Core Idea:

The core idea of this project is to design and develop a comprehensive data pipeline integration and music recommendation system using Docker and Amazon Web Services (AWS). The system leverages the Spotify Web API to extract music data, including playlists and tracks, and integrates it with AWS services to build a scalable and efficient data pipeline.

Handling large volumes of music data poses a challenge in the music industry, where traditional data pipelines are inefficient at managing data scale and complexity. This project addresses this issue by implementing a Dockerized data pipeline that ensures portability and consistency, while AWS services provide the infrastructure to handle large-scale data processing.

The system consists of four key components: data ingestion, data transformation, data loading, and music recommendation. Data is extracted using Dockerized containers and stored in Amazon S3. Data transformation is handled by Docker and AWS Glue, and the processed data is made available for analysis using Amazon Athena. A collaborative filtering algorithm generates personalized music recommendations based on user listening history. The technical approach integrates Docker with AWS services, including ECS, S3, Glue, and Athena, to create a highly scalable and efficient data pipeline. The innovative aspect lies in Docker's role in maintaining environment consistency and the use of collaborative filtering for generating music recommendations.

Project Setup and Configuration

Repository Structure

```
Spotify_E2E_ETL_DataEngineering_project/
├─ data/
├─ docs/
├─ src/
│   ├─ extract/
│   ├─ transform/
│   └─ load/
├─ utils/
├─ tests/
├─ .gitignore
├─ README.md
├─ requirements.txt
└─ setup.py
```

Prerequisites

- Docker
- jsons==1.6.3
- boto3==1.35.29
- datetime==5.5
- Python-IO==0.3
- pandas==2.2.3
- spotipy==2.24.0

Installation Guide

1. **Clone the repository:**
git clone https://github.com/sudarshanp1/Spotify_E2E_ETL_DataEngineering_project.git
cd Spotify_E2E_ETL_DataEngineering_project
2. **Install the required packages:**
pip install -r requirements.txt
3. **Set up environment variables:**
Create a .env file in the root directory with your Spotify API credentials and database connection details.

Data Pipeline Overview

Data Sources

The primary data source for this project is the Spotify Web API, which provides various endpoints to fetch data related to tracks, artists, albums, and playlists.

Data Ingestion

Data is ingested using API calls to Spotify's endpoints. The ingestion process is handled by the scripts in the src/extract directory.

Data Transformation

Data transformation involves cleaning, normalizing, and enriching the raw data. This is done using the scripts in the src/transform directory.

Data Storage

The transformed data is stored in a PostgreSQL database. The schema design and database setup are covered in the next section

Data Modeling

Schema Design

The database schema is designed to efficiently store and query Spotify data. It includes tables for tracks, artists, albums, playlists, and users.

Database Setup

1. Create a PostgreSQL database.
2. Run the schema SQL scripts located in the docs/sql directory to set up the tables.

ETL Processes

Extract

The **data extraction** process involves making API calls to the Spotify Web API to retrieve relevant data. This is implemented using the **requests** library within the src/extract module, which is containerized using Docker for portability. The extracted data includes metadata related to tracks, artists, albums, and playlists and is stored in its raw form in an Amazon S3 bucket.

Transform

Data transformation focuses on cleaning and enriching the extracted data. The transformation process includes:

- **Handling missing values:** Filling in missing data points or removing incomplete records.
- **Normalizing text fields:** Ensuring consistent formatting across text-based fields (e.g., song and artist names).
- **Enriching data:** Adding additional metadata, such as genre classification or popularity metrics. This process is handled by Dockerized scripts in the src/transform directory and AWS Glue for more complex transformations.

Load

The **load** process inserts the transformed data into the PostgreSQL database. This step is managed by the src/load module, which establishes a connection to the PostgreSQL instance and efficiently loads the cleaned data into the appropriate schema. The transformed data is now ready for querying and analysis through tools like AWS Athena or PostgreSQL queries.

Implementation Details

Docker Integration

- Docker ensures portability across environments for consistent data processing.

Code Walkthrough

- `src/extract/api_extractor.py`: Handles API calls to Spotify.
- `src/transform/data_cleaner.py`: Contains functions for data cleaning and normalization.
- `src/load/db_loader.py`: Manages database connections and data insertion.

Key Functions and Classes

- `APIExtractor`: Class responsible for making API requests and handling responses.
- `DataCleaner`: Class for cleaning and transforming raw data.
- `DBLoader`: Class for loading data into the database.

4. Testing and Validation

Unit Testing

Unit tests are written using the unittest framework and are located in the tests directory.

Integration Testing

Integration tests ensure that the different components of the ETL pipeline work together seamlessly.

5. Deployment

Deployment Strategies

The project uses Docker for containerization, making it easy to deploy the ETL pipeline across different environments.

CI/CD Pipeline

A CI/CD pipeline is set up using GitHub Actions to automate testing and deployment.

Technical Implementation

Setting Up AWS Services

Amazon S3:

Create an S3 bucket to store raw and transformed data extracted from the Spotify API.

AWS Lambda:

Create a Lambda function to handle the extraction of data from the Spotify API and load it into the S3 bucket.

AWS Glue Crawler:

Set up a Glue Crawler to crawl the S3 bucket and create a Glue Data Catalog.

AWS Glue Data Catalog:

Use the Glue Data Catalog to manage and maintain metadata about the data stored in S3.

Amazon Athena:

Set up Athena to query the data stored in the Glue Data Catalog.

CloudWatch:

Set up CloudWatch to monitor the AWS resources and Lambda functions.

Spotify API Data Extraction in Docker:

Containerize the Lambda function for extracting data from the Spotify API using Docker to ensure consistency across development and production environments.

Dockerfile:

- Create a Dockerfile to define the environment and dependencies required for running the Spotify API extraction logic.

Local Development:

- Use Docker containers to develop and test the Lambda function locally before deploying it to AWS.

Amazon ECS (Elastic Container Service):

- For larger-scale or batch extractions, deploy the Docker container to Amazon ECS, which allows for scalable container orchestration.

Detailed Steps for ImplementationAWS Setup

- Create an S3 bucket.
- Set up IAM roles for Lambda and Glue with necessary permissions.

Spotify API Integration

- Register an application on the Spotify Developer Dashboard to get the client ID and client secret.
- Test API calls using Postman or a similar tool to ensure proper data retrieval.

Lambda Function Deployment

- Use the AWS Management Console or AWS CLI to create and deploy the Lambda function.
- Set up a CloudWatch event to trigger the Lambda function periodically.

Glue Crawler and Jobs

- Create a Glue Crawler in the AWS Management Console and configure it to crawl the S3 bucket.
- Develop and test the Glue ETL job in the Glue Studio or using PySpark scripts.
- Schedule the Glue job to run after the data extraction Lambda function.


Data Analysis with Athena

- Create an Athena database and table to query the transformed data.
- Use SQL queries in Athena to analyze the data and create reports.

Extract Data from Spotify API

- Spotify Developer Account:
- Sign up for a Spotify Developer account and create an app to get API credentials.
- Lambda Function for Data Extraction:
- Write a Lambda function in Python to extract data from the Spotify API and load it into the S3 bucket.

This method involves using AWS Lambda to extract music data from the Spotify Web API and store it in an Amazon S3 bucket.



1. Create spotify-data-extractor directory
2. Create Dockerfile file

```
FROM python:3.12
WORKDIR /usr/local/app

# Install the application dependencies
COPY requirements.txt ./
RUN pip install --no-cache-dir -r requirements.txt

# Copy in the source code
COPY . .
EXPOSE 80

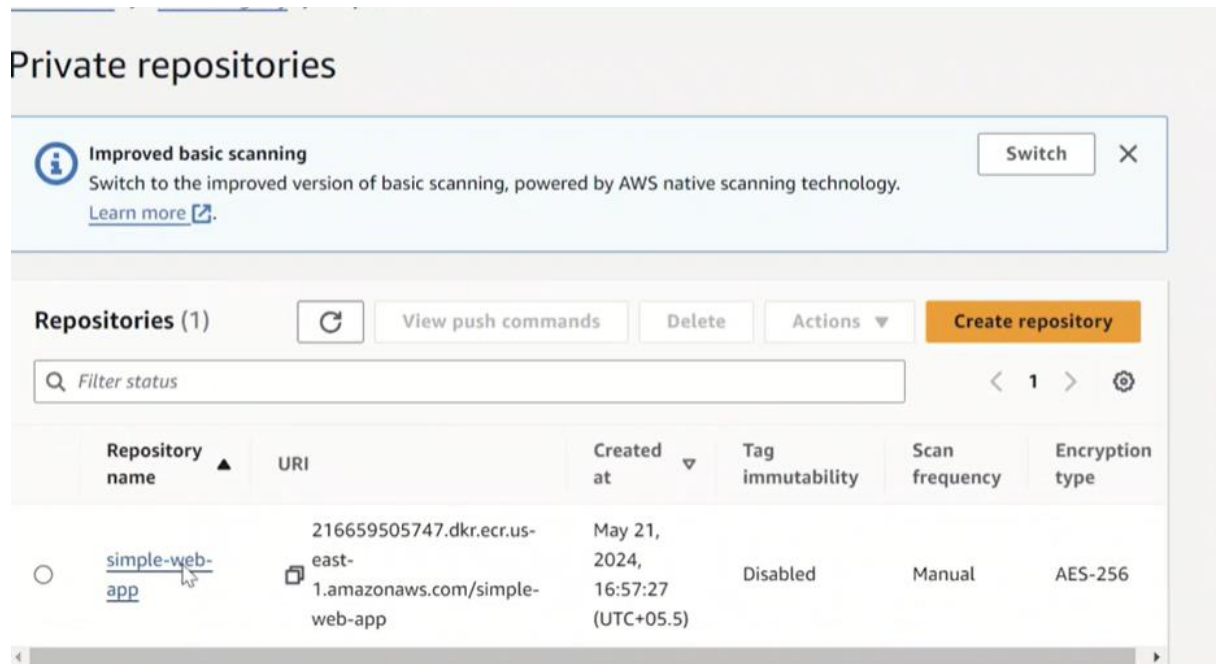
CMD ["python", "spotify_api_data_extract.py"]
```

3. Create requirement file

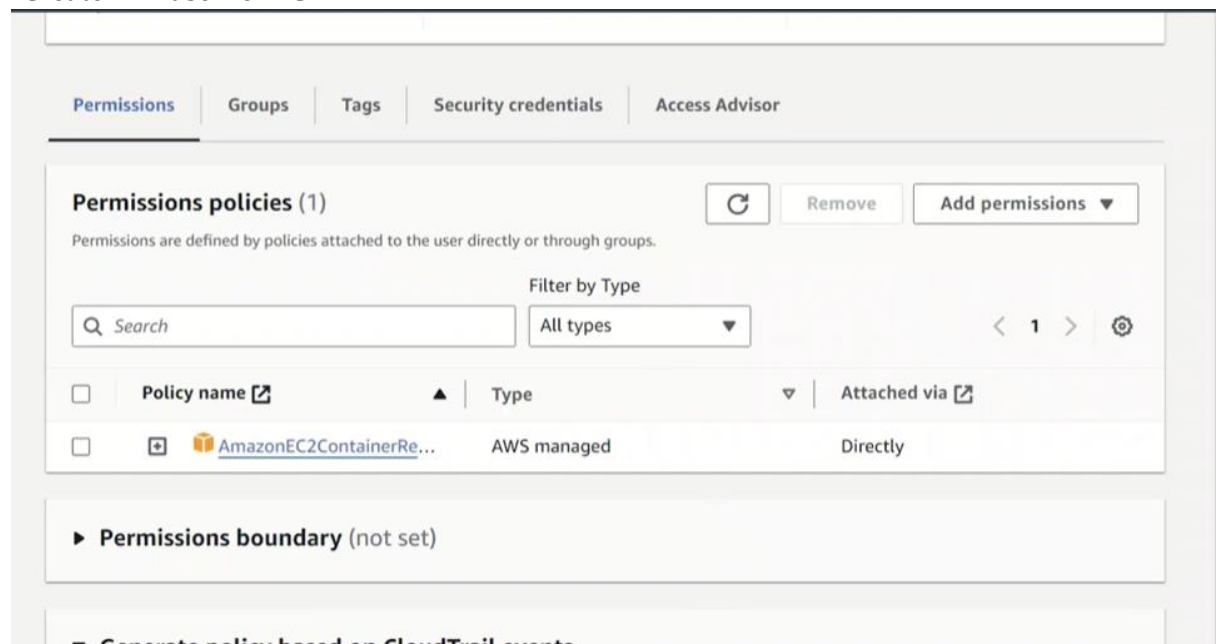
```
jsons==1.6.3
boto3==1.35.29
datetime==5.5
Python-IO==0.3
pandas==2.2.3
spotipy==2.24.0
```

Setup AWS CLI

4. Create ECR



5. Create IAM user for ECR



6. Create a access key

7. Authenticate docker to ECR

```
AWS Access Key ID [*****7EGF]: AKIATE4PJWJJVGORQBNX
AWS Secret Access Key [*****iadM]: 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
Default region name [us-east-1]:
Default output format [None]:
```

8. Create docker image

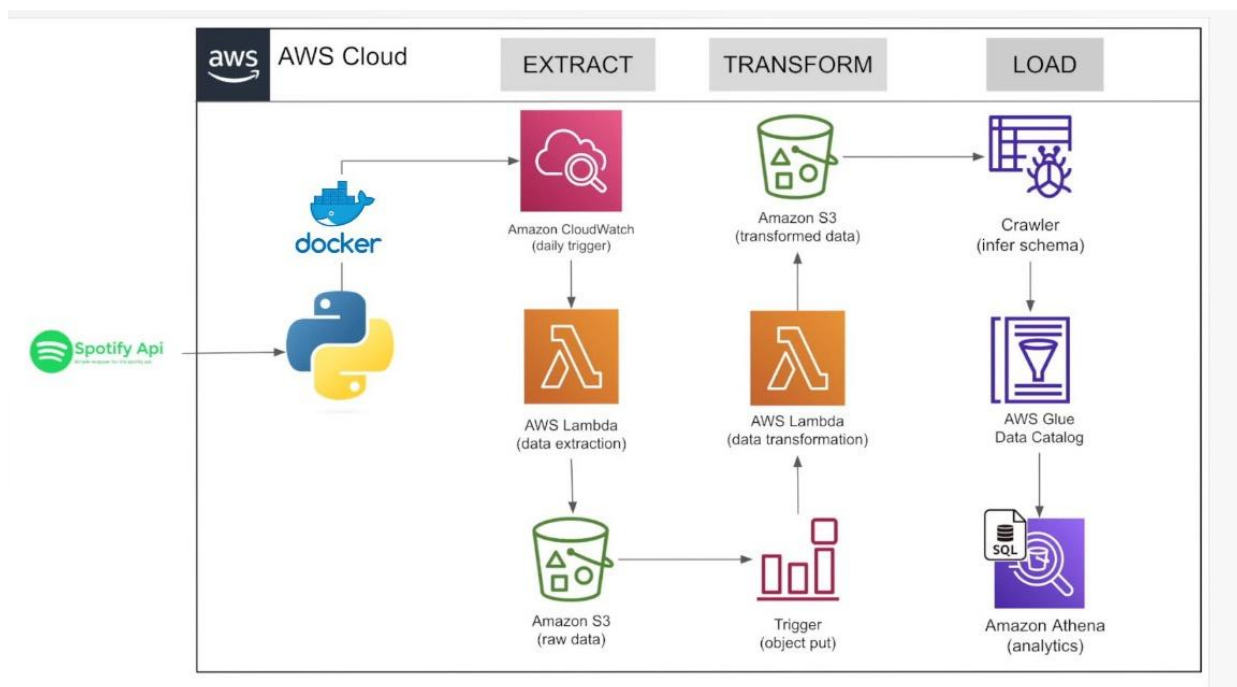
```
PS C:\Users\abhi8\spotify-data-extractor> docker build -t spotify-data-extractor .
[+] Building 1.0s (10/10) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 301B
=> [internal] load metadata for docker.io/library/python:3.12
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load build context
=> => transferring context: 116B
=> [1/5] FROM docker.io/library/python:3.12@sha256:14f073695854184b65a82808ea2baa352d49c1a86442f0
=> CACHED [2/5] WORKDIR /usr/local/app
=> CACHED [3/5] COPY requirements.txt ./
=> CACHED [4/5] RUN pip install --no-cache-dir -r requirements.txt
=> CACHED [5/5] COPY . .
=> exporting to image
=> => exporting layers
=> => writing image sha256:821f4d0d9d397876c00d60a3c07385ea68dc945af48e08f9f5dd0851f371b637
=> => naming to docker.io/library/spotify-data-extractor

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/0xa46tcqkhlsljlt
```

9. Tag the docker image

10. Push the docker image

```
the push refers to repository [216659505747.dkr.ecr.us-east-1.amazonaws.com/simple-web-app]
593ee8fbd7b: Pushed
e495f7b0b7d: Pushed
c70f446fbe2: Pushed
be225e16e44: Pushed
d04ead9b400: Pushed
f5598fef05f: Pushed
fbd5a835e5e: Pushed
5061be64847: Pushed
4fc045c9e3a: Pushed
```



sudarshanp1 / Spotify_E2E_ETL_DataEngineering_project

🔍 Type / to search

<> Code

Issues

Pull requests

Actions

Projects

Wiki

Security

Insights

Settings

Files

main

Go to file

01_README.md

ETL_Pipeline_Spotify1.ipynb

Spotify_Data_Pipeline.jpeg

Spotify_api_data_extract.png

spotify_api_data_extract.py

spotify_transformation_load_func...

spotify_transformation_load_func...

Spotify_E2E_ETL_DataEngineering_project / spotify_transformation_load_function.py

sudarshanp1 Create spotify_transformation_load_function.py 3f33bdd

Code Blame 113 lines (94 loc) · 4.35 KB Code 55% faster with GitHub Copilot

```
1
2 #!/usr/bin/env python
3 # coding: utf-8
4
5 # In[ ]:
6
7
8 import json
9 import boto3
10 from datetime import datetime
11 from io import StringIO
12 import pandas as pd
13
14 def album(data):
15     album_list = []
16     for row in data['items']:
17         album_id = row['track']['album']['id']
18         album_name = row['track']['album']['name']
19         album_release_date = row['track']['album']['release_date']
20         album_total_tracks = row['track']['album']['total_tracks']
21         album_url = row['track']['album']['external_urls']['spotify']
22         album_element = {'album_id':album_id,'name':album_name,'release_date':album_release_date,
23                         'total_tracks':album_total_tracks,'url':album_url}
24         album_list.append(album_element)
25     return album_list
26
27 def artist(data):
28     artist_list = []
29     for row in data['items']:
30         for key, value in row.items():
31             if key == "track":
32                 for artist in value['artists']:
33                     artist_dict = {'artist_id':artist['id'], 'artist_name':artist['name'], 'external_url': artist['href']}
34                     artist_list.append(artist_dict)
35     return artist_list
```

sudarshanp1 / Spotify_E2E_ETL_DataEngineering_project

🔍 Type / to search

<> Code

Issues

Pull requests

Actions

Projects

Wiki

Security

Insights

Settings

Files

main

Go to file

01_README.md

ETL_Pipeline_Spotify1.ipynb

Spotify_Data_Pipeline.jpeg

Spotify_api_data_extract.png

spotify_api_data_extract.py

spotify_transformation_load_func...

spotify_transformation_load_func...

Spotify_E2E_ETL_DataEngineering_project / spotify_api_data_extract.py

sudarshanp1 spotify_api_data_extract.py b04aa9

Code Blame 37 lines (25 loc) · 1.07 KB Code 55% faster with GitHub Copilot

```
1
2 #!/usr/bin/env python
3 # coding: utf-8
4
5 # In[ ]:
6
7
8 import json
9 import os
10 import spotipy
11 from spotipy.oauth2 import SpotifyClientCredentials
12 import boto3
13 from datetime import datetime
14
15 def lambda_handler(event, context):
16
17     client_id = os.environ.get('client_id')
18     client_secret = os.environ.get('client_secret')
19
20     client_credentials_manager = SpotifyClientCredentials(client_id = client_id, client_secret = client_secret)
21     sp = spotipy.Spotify(client_credentials_manager = client_credentials_manager)
22     playlists = sp.user_playlists('spotify')
23
24     playlist_link = "https://open.spotify.com/playlist/3719d02EVXtNG2KDCfK0F?si=1333723a6eff4b7f&nd=1&lsi=1586c2ede2dc4823"
25     playlist_URI = playlist_link.split("/")[-1].split('?')[0]
26
27     spotify_data = sp.playlist_tracks(playlist_URI)
28
29     client = boto3.client('s3')
30
31     filename = "spotify_raw_" + str(datetime.now()) + ".json"
32
33     client.put_object(
34         Bucket = "spotify-etl-project-kushank",
35         Key = "raw_data/to_processed/" + filename,
36         Body = json.dumps(spotify_data)
37     )
```


The image displays two overlapping screenshots from the AWS ecosystem. The top screenshot shows the AWS IDE editor with a Python script named `index.mjs`. The script uses `json`, `boto3`, `pandas`, `StringIO`, and `datetime` modules. It defines three functions: `album(data)`, `artists(data)`, and `songs(data)`, which process data from a JSON input. The bottom screenshot shows the Amazon S3 console for the bucket `spotify-etl-project-ak`. Under the 'Objects' tab, three items are listed: `raw_data/` (Folder), `spotify.zip` (zip), and `transformed_data/` (Folder). A mouse cursor is pointing at the `transformed_data/` folder.

```
1 import json
2 import boto3
3 import pandas as pd
4 from io import StringIO
5 from datetime import datetime
6
7 def album(data):
8     album_list = []
9     for row in data['items']:
10
11         album_id = row['track']['album']['id']
12         album_name = row['track']['album']['name']
13         album_release_date = row['track']['album']['release_date']
14         album_total_tracks = row['track']['album']['total_tracks']
15         album_url = row['track']['album']['external_urls']['spotify'] #now we need to create an album element out of all
16         album_element = {'album_id': album_id, 'album_name': album_name, 'album_release_date': album_release_date, 'album_total_tracks': album_total_tracks, 'album_url': album_url}
17         album_list.append(album_element)
18     return album_list
19
20
21 def artists(data):
22     artists_list = []
23     for row in data['items']:
24         song_id = row['track']['id']
25         for artists in row['track']['artists']:
26             artists_element = {
27                 'artists_id': artists['id'], 'song_id': song_id, 'artist_name': artists['name'], 'external_url': artists['href']
28             }
29             artists_list.append(artists_element)
30
31
32 def songs(data):
33     song_list = []
34     for row in data['items']:
35         song_id = row['track']['id']
36         song_name = row['track']['name']
37         song_duration = row['track']['duration_ms']
38         song_element = {'song_id': song_id, 'song_name': song_name, 'song_duration': song_duration}
```

Amazon S3 > Buckets > spotify-etl-project-ak

spotify-etl-project-ak [Info](#)

[Objects](#) | [Properties](#) | [Permissions](#) | [Metrics](#) | [Management](#) | [Access Points](#)

Objects (3) [Info](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket.

Find objects by prefix

<input type="checkbox"/>	Name	Type
<input type="checkbox"/>	<code>raw_data/</code>	Folder
<input type="checkbox"/>	<code>spotify.zip</code>	zip
<input type="checkbox"/>	<code>transformed_data/</code>	Folder

4.1 Data Transformation using AWS lambda

Schedule a CloudWatch event to trigger a specific Lambda function at a recurring interval, with the following configuration:

Schedule Expression: rate(1 hour) to execute the function every hour

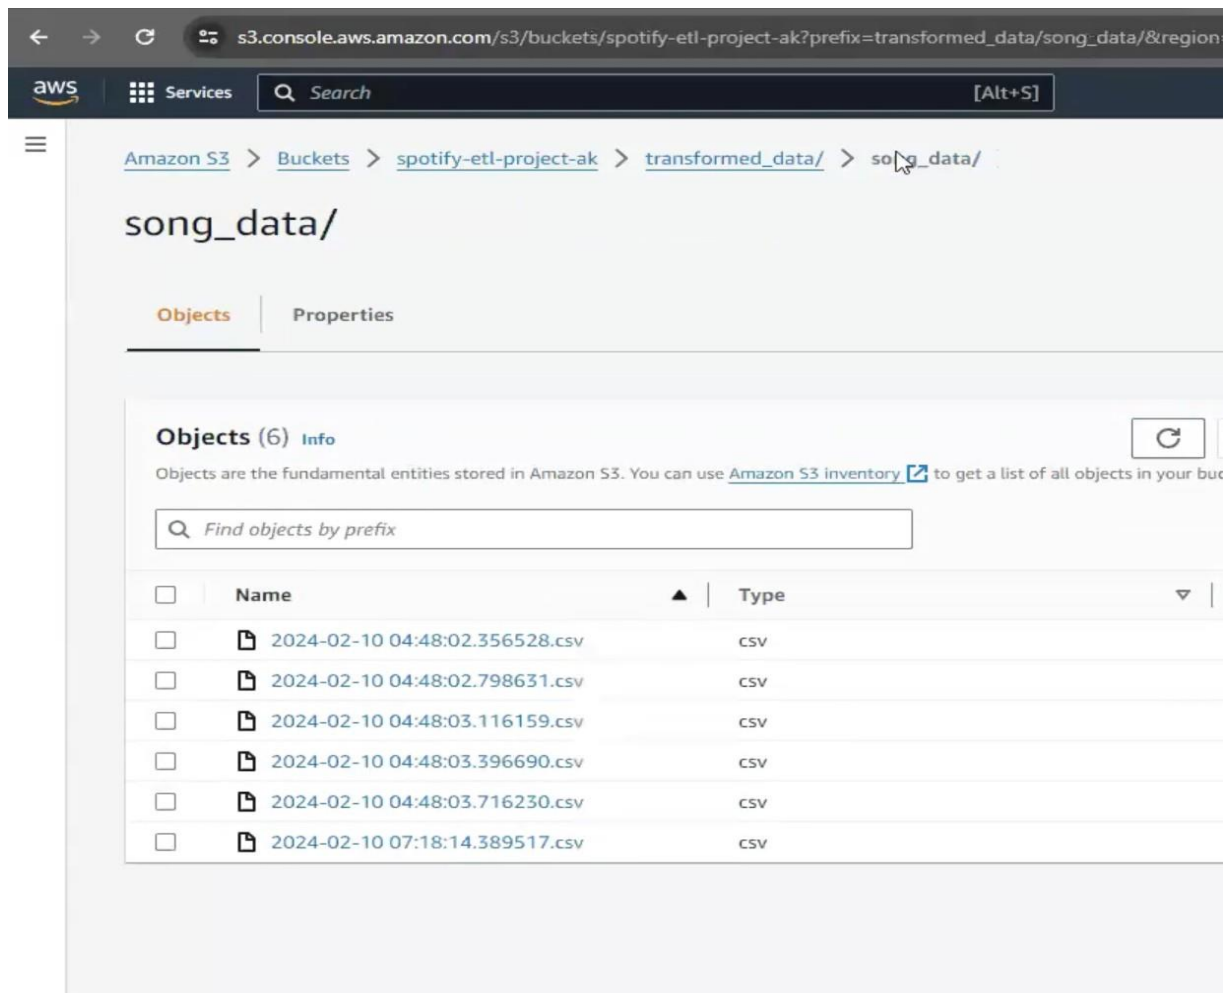
Trigger Type: CloudWatch Events - Schedule

Rule Description: Automatically run the Lambda function at a specified frequency (e.g., daily or weekly)

This setup will ensure the Lambda function is invoked at the desired rate, providing a consistent and reliable execution schedule.

Transform Data using AWS Glue

- Create Glue Crawler:
- Set up a Glue Crawler to crawl the S3 bucket and create a Data Catalog.
- Configure the crawler to run periodically to keep the catalog updated.
- Transform Data with Glue Jobs:
- Create a Glue Job to transform the raw data into a structured format.



us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/spotify_data_transform/tab=configure

aws Services Search [Alt+S]

Function overview Info

Export to Application Composer

Diagram Template

spotify_data_transform

Layers (1)

S3

+ Add trigger

+ Add destination

Description

-

Last modified 3 hours ago

Function ARN

arn:aws:lambda:us-east-1:::tion:spotify_data_transform

Function URL Info

-

Code Test Monitor Configuration Aliases Versions

General configuration

Triggers

Triggers (1) Info

Find triggers

Fix errors Edit Delete

us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/spotify-extract-api-data/tab=configure

aws Services Search [Alt+S]

Function overview Info

Export to Application Composer

Diagram Template

spotify-extract-api-data

Layers (1)

EventBridge (CloudWatch Events)

+ Add trigger

+ Add destination

Description

-

Last modified 12 minutes ago

Function ARN

arn:aws:lambda:us-east-1:355tion:spotify-extract-api-data

Function URL Info

-

Code Test Monitor Configuration Aliases Versions

General configuration

Triggers

Triggers (1) Info

Find triggers

Fix errors Edit Delete

Load Data into AWS Data Store

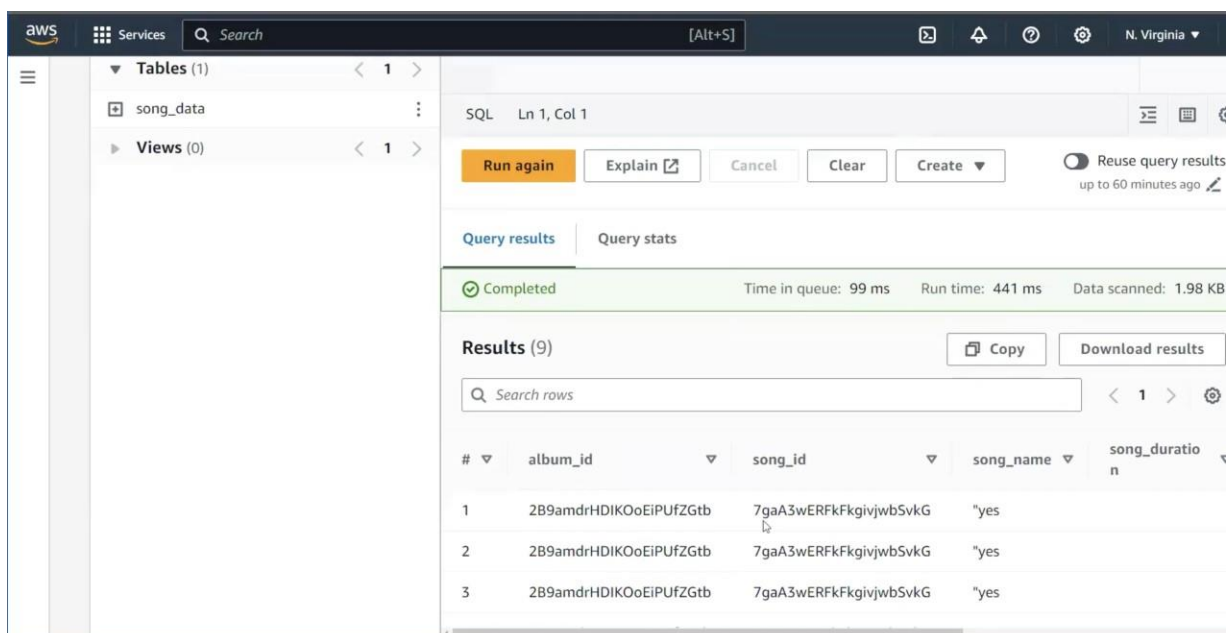
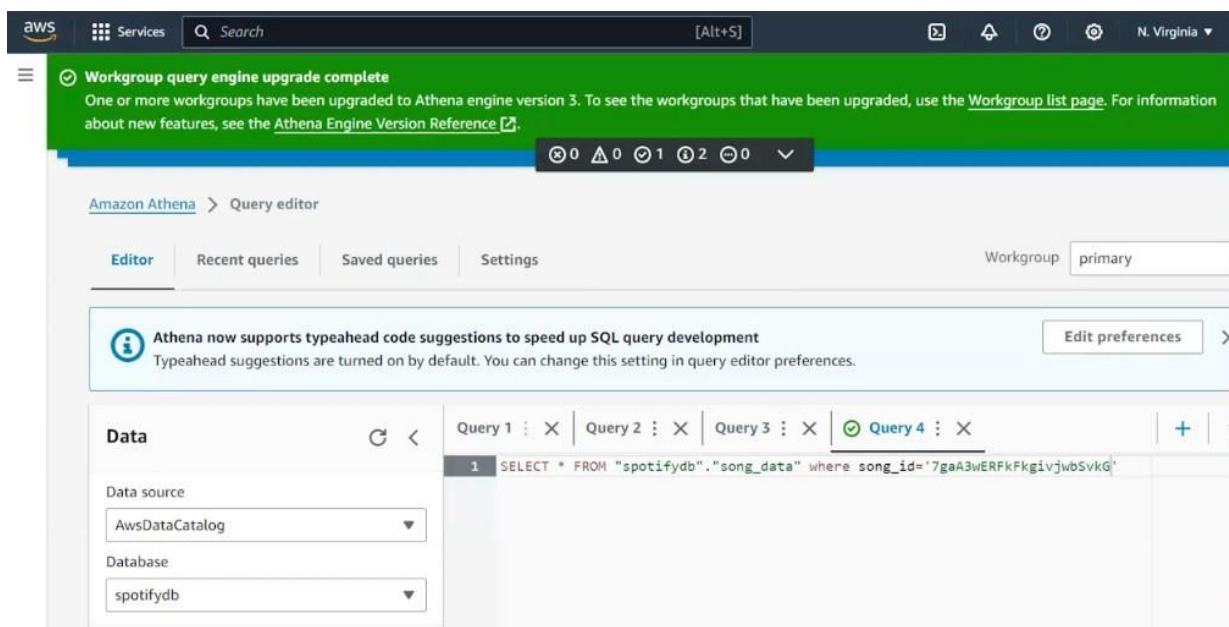
- Run Glue Crawler:
- Run the Glue Crawler to update the Data Catalog with the transformed data.
- Query Data with Athena:
- Use Athena to query the transformed data stored in S3.

Data Preparation

We successfully created tables in Amazon Athena to match the data stored in Amazon S3. This included setting up an external table for Song Data, which includes important details like song ID and more.

User Behavior Analysis

We analyzed user behavior and found some interesting patterns. We saw which songs were played the most and which ones were skipped often. We also looked at when users listened to music throughout the day, which will help us make better recommendations in the future.



The screenshot shows the AWS Athena console interface. On the left, there's a sidebar with 'Tables (1)' and 'Views (0)'. The main area displays a query result for a table named 'song_data'. The query is completed, and the results are shown in a table with 9 rows. The columns are 'album_id', 'song_id', 'song_name', and 'song_duration'. The first three rows are visible, showing song IDs and names.

#	album_id	song_id	song_name	song_duration
1	2B9amdrHDIKooEiPUfZGtb	7gaA3wERFkFkgivjwbSvkG	*yes	
2	2B9amdrHDIKooEiPUfZGtb	7gaA3wERFkFkgivjwbSvkG	*yes	
3	2B9amdrHDIKooEiPUfZGtb	7gaA3wERFkFkgivjwbSvkG	*yes	

Monitoring and Reporting

CloudWatch Alarms:

Set up CloudWatch alarms to monitor Lambda function errors, Glue job failures, and other critical metrics.

CloudWatch Logs:

Enable CloudWatch logs for Lambda functions and Glue jobs to track and troubleshoot issues.

Reporting:

Generate reports based on the queries run in Athena and visualize the data using tools like Amazon QuickSight or other BI tools.

Collaborative Filtering for Music Recommendation

This method involves using collaborative filtering to generate personalized song recommendations for users based on their listening history and preferences.

System Architecture and Deployment on AWS

This method involves deploying the system architecture on AWS, using AWS services to build a scalable and efficient data pipeline.

Conclusion

Achievements

- Successfully designed and developed an end-to-end ETL pipeline to extract, transform, and load Spotify data.
- Ensured data quality, accuracy, and integrity throughout all stages of the pipeline, from extraction to storage.
- Delivered meaningful insights and actionable analytics from the processed Spotify data, demonstrating the system's capability to handle large-scale music data efficiently.

Future Work

- Enhance data transformation processes with more advanced analytics.
- Implement real-time data processing capabilities.
- Expand the data model to include additional Spotify data.

References

1. Spotify Web API Documentation: <https://developer.spotify.com/documentation/web-api/>
2. PostgreSQL Documentation: <https://www.postgresql.org/docs/>

Discussion: Data Pipeline Integration & Music Recommendation System on AWS (with Docker)

The data pipeline integration and music recommendation system developed in this project leverages both AWS services and Docker containerization to provide a scalable, efficient, and flexible solution for handling large volumes of music data. Below is a discussion of the advantages and disadvantages of this approach:

Advantages:

- **Scalability:** The system is designed to handle large volumes of music data and can scale to meet the needs of a growing user base. Using Docker containerization ensures that new instances can be deployed quickly and consistently, while AWS services such as ECS, Fargate, and S3 ensure scalability in terms of both data processing and infrastructure.
- **Efficiency:** The use of AWS services, including Dockerized containers running on ECS, S3 for storage, Glue for data transformation, and Athena for analysis, enables efficient data processing. Docker containers further improve efficiency by isolating workloads and allowing smooth deployments and updates.
- **Containerization Flexibility:** Docker containers allow the application to be easily portable across different environments (e.g., local, development, and production), ensuring consistency and simplifying the development workflow. It also enables isolated environments for different services (e.g., data extraction, transformation) within the pipeline.
- **Personalization:** The music recommendation system uses collaborative filtering to generate personalized song recommendations based on user listening history and preferences. This provides a tailored experience for each user.
- **Cost-Effectiveness:** By utilizing AWS services like ECS with Fargate and S3, the system remains cost-efficient by only paying for the resources used. Docker containers also reduce overhead by optimizing resource utilization.
- **Flexibility:** Docker containerization and AWS services enable the system to be easily modified and extended to incorporate new features and functionality. Containers allow for quick iterations and updates to different components of the system.

Disadvantages:

- **Complexity:** Integrating Docker with AWS services adds an additional layer of complexity to the system. Managing Docker images, orchestrating containers, and deploying across different services requires significant technical expertise.
- **Dependence on AWS:** The system relies heavily on AWS services. Any downtime or service disruptions in AWS could impact system performance or availability. This dependence is amplified when using both AWS and Docker, as orchestration is tightly coupled with AWS infrastructure.
- **Data Quality:** The quality of the music data extracted from the Spotify Web API can directly affect the accuracy of the recommendation system. Issues with data inconsistencies or missing information may degrade the quality of recommendations.
- **Cold Start Problem:** Like many recommendation systems, the cold start problem occurs when new users do not have sufficient listening history, making it difficult to generate personalized recommendations.
- **Spotify API Limitations:** The Spotify Web API imposes limits on the amount of data that can be extracted, which can impact the scalability and comprehensiveness of the data used in the system. Rate limits and data availability restrictions may hinder large-scale operations.

Future Improvements and Enhancements:

- **Improving Data Quality:** Implement automated data validation and quality checks to ensure that the data extracted from the Spotify API is accurate, complete, and consistent. This would improve the performance of the recommendation engine.
- **Addressing the Cold Start Problem:** Develop strategies to mitigate the cold start issue, such as incorporating content-based filtering or hybrid recommendation techniques to enhance personalization for new users.
- **Enhancing Personalization:** Leverage more advanced personalization techniques, such as natural language processing (NLP) or deep learning algorithms, to provide better song recommendations based on user behavior and preferences.
- **Expanding to Other Music Platforms:** Integrate data from other music streaming platforms like Apple Music or Google Play Music to provide more comprehensive insights and recommendations. This would also make the system platform-agnostic, expanding its potential user base.
- **Optimizing System Performance:** Improve the performance of the system by further optimizing Docker containers and AWS services. For example, reducing latency, improving query execution times in Athena, and optimizing data transformations in Glue can enhance overall efficiency.
- **Implementing Real-Time Data Processing:** Integrating real-time data processing capabilities (e.g., AWS Kinesis with Docker for streaming) would allow the system to provide up-to-the-minute recommendations and analytics, improving user engagement.

By continuously improving both the Docker containerization and AWS-based architecture, the system can scale efficiently while providing users with increasingly accurate and personalized music recommendations.

Data Pipeline Integration & Music Recommendation System on AWS

The data pipeline integration and music recommendation system on AWS developed in this project has several advantages and disadvantages.

Advantages:

Scalability: The system is designed to handle large volumes of music data and can scale to meet the needs of a growing user base.

Efficiency: The use of AWS services, including Lambda, S3, Glue, and Redshift, enables efficient data processing and analysis.

Personalization: The music recommendation system uses collaborative filtering to generate personalized song recommendations for users based on their listening history and preferences.

Cost-Effective: The use of AWS services provides a cost-effective solution for building and deploying the system.

Flexibility: The system can be easily modified and extended to incorporate new features and functionality.

Disadvantages:

Complexity: The system is complex and requires significant technical expertise to design, develop, and deploy.

Dependence on AWS: The system is dependent on AWS services, which can be a risk if AWS experiences downtime or service disruptions.

Data Quality: The quality of the music data extracted from the Spotify Web API can affect the accuracy of the music recommendation system.

Cold Start Problem: The system may experience a cold start problem, where new users do not have enough listening history to generate accurate music recommendations.

Spotify API Limitations: The Spotify Web API has limitations on the amount of data that can be extracted, which can affect the scalability of the system.

Future Improvements and Enhancements

There are several areas for future improvements and enhancements, including:

Improving Data Quality: Implementing data quality checks and validation to ensure the accuracy of the music data extracted from the Spotify Web API.

Addressing the Cold Start Problem: Developing strategies to address the cold start problem, such as using content-based filtering or hybrid approaches.

Enhancing Personalization: Developing more advanced personalization techniques, such as using natural language processing or deep learning algorithms.

Expanding to Other Music Platforms: Expanding the system to incorporate music data from other platforms, such as Apple Music or Google Play Music.

Improving System Performance: Optimizing system performance to improve response times and reduce latency.